P536

Advanced Operating Systems

Computer Science Department
Indiana University
Bloomington, IN 47405

Home · **About** · **Schedule** · **Assignments** · **Resources**   **May 29, 2013**

Main Menu

◇ Home
◇ Contact
◇ About
◇ Schedule
◇ Assignments
◇ Resources
◇ Fun

The Tao

The Magician of the Ivory Tower brought his latest invention for the master programmer to examine. The magician wheeled a large black box into the master's office while the master waited in silence.

"This is an integrated, distributed, general-purpose workstation," began the magician, "ergonomically designed with a proprietary operating system, sixth generation languages, and multiple state of the art user interfaces. It took my assistants several hundred man years to construct. Is it not amazing?"

The master raised his eyebrows slightly. "It is indeed amazing," he said.

"Corporate Headquarters has commanded," continued the magician, "that

Programming Assignment 3: Virtual Memory

Assigned	Oct 3, 2012
Design Due	Oct 12, 2012
Due	Nov 7, 2012
Note	This problem set is to be done in groups of two.

Introduction

You have improved OS/161 to the point that you can now run user processes. However, there are a number of shortcomings in the current system. A process's size is limited by the number of TLB entries (i.e., 64 pages). In this assignment we will adapt OS/161 to take full advantage of the simulated hardware by implementing management of the MIP software-managed Translation Lookaside Buffer (TLB). You will write the code to manage this TLB. You will also write the code to implement paging -- the mechanism by which memory pages of an active process can be sent to disk when memory is needed, and restored to memory when required by the program. This permits many processes to share limited physical memory while providing each process with the abstraction of a very large virtual memory.

Structure of the TLB entries

In the System/161 machine, each TLB entry includes a 20-bit virtual page number and a 20-bit physical page number as well as the following five fields:

- `global`: 1 bit; if set, ignore the `pid` bits in the TLB.
- `valid`: 1 bit; set if the TLB entry contains a valid translation.
- `dirty`: 1 bit; enables writing to the page referenced by the entry; if this bit is 0, the page is only accessible for reading.
- `nocache`: 1 bit; unused in System/161. In a real processor, indicates that the hardware cache will be disabled when accessing this page.
- `pid`: 6 bits; a context or address space ID that can be used to allow entries to remain in the TLB after a context switch.

All these bits/values are maintained by the operating system. When the `valid` bit is set, the TLB entry contains a valid translation. This implies that the virtual page is present in physical memory. A **TLB miss** occurs when no TLB entry can

everyone use this
workstation as a
platform for new
programs. Do you agree
to this?"

"Certainly," replied the
master, "I will have it
transported to the data
center immediately!"
And the magician
returned to his tower,
well pleased.

Several days later, a
novice wandered into
the office of the master
programmer and said, "I
cannot find the listing
for my new program. Do
you know where it
might be?"

"Yes," replied the
master, "the listings are
stacked on the platform
in the data center."

— Geoffrey James, *The
Tao of Programming*

be found with a matching virtual page and address space ID (unless the global
bit is set in which case the address space ID is ignored) with the valid bit set.

Paging

The operating system creates the illusion of unlimited memory by using physical
memory as a cache of **virtual pages**. Paging relaxes the requirement that all
the pages in a process's virtual address space must be in physical memory.
Instead, we allow a process to have pages either on disk or in memory. When
the process issues an access to a page that is on disk, a **page fault** occurs. The
operating system must retrieve the page from disk and bring it into memory.
Pages with valid TLB entries are always in physical memory. This means that a
reference to a page on disk will always generate a TLB fault. At the time of a
TLB fault, the hardware generates a TLB exception, trapping to the operating
system. The operating system then checks its own page table to locate the
virtual page requested. If that page is currently in memory but wasn't mapped
by the TLB, then all we need to do is update the TLB. However, the page might
be on disk. If this is the case, the operating system must:

1. Allocate a place in physical memory to store the page;
2. Read the page from disk,
3. Update the page table entry with the new virtual-to-physical address
   translation;
4. Update the TLB to contain the new translation; and
5. Resume execution of the user program.

Notice that when the operating system selects a location in physical memory in
which to place the new page, the space may already be occupied. In this case,
the operating system must **evict** that other page from memory. If the page has
been modified or does not currently have a copy on disk, then the old page must
first be written to disk before the physical page can be reallocated. If the old
page has not been modified and already has a copy on disk, then the write to
disk can be avoided. The appropriate page table entry must be updated to
reflect the fact that the page is no longer in memory.

As with any caching system, performance of your virtual memory system
depends on the policy used to decide which things are kept in memory and
which are evicted. On a page fault, the kernel must decide which page to
replace. Ideally, it will evict a page that will not be needed soon. Many systems
(such as UNIX) avoid the delay of synchronously writing memory pages to disk
on a page fault by writing modified pages to disk in advance, so that subsequent
page faults can be completed more quickly.

Your Mission Your mission is to:

- Implement the code that services TLB faults.
- Add paging to your operating system.
- Tune your TLB and memory replacement algorithms.
- Add the `sbrk()` system call, so that the `malloc()` we give you works

Setup

Before you begin this assignment, tag your repository as `asst3-begin`.

Consult the ASST3 config file and notice that the `arch/mips/mips/dumbvm.c` file

will be omitted from your kernel. You will undoubtedly need to add new files to the system for this assignment, e.g., `kern/vm/vm.c` or `kern/arch/mips/mips/mipsvm.c`. Be sure to update the file `kern/conf/conf.kern`, or, for machine-dependent files, `kern/arch/mips/conf/conf.arch`, to include any new files that you create. Take care to place files in the "correct" place, separating machine-dependent components from machine-independent components.

You should also now restrict your physical memory to 512 KB by editing the `ramsize` line in your `sys161.conf` file.

Now, config an ASST3 kernel, run `make depend`, and build it. You are now ready to begin assignment 3; tag your repository `asst3-begin`.

Design Questions

Place answers to the following questions in `questions.txt`.

Problem 1 (5 points)

Assuming that a user program just accessed a piece of data at (virtual) address X, describe the conditions under which each of the following can arise. If the situation cannot happen, answer "impossible" and explain why it cannot occur.

- TLB miss, page fault
- TLB miss, no page fault
- TLB hit, page fault
- TLB hit, no page fault

Problem 2 (2 points)

A friend of yours who foolishly decided not to take 161, but who likes OS/161, implemented a TLB that has room for only one entry, and experienced a bug that caused a user-level instruction to generate a TLB fault infinitely - the instruction never completed executing! Explain how this could happen. (Note that after OS/161 handles an exception, it restarts the instruction that caused the exception.)

Problem 3 (3 points)

How many memory-related exceptions (i.e., hardware exceptions and other software exceptional conditions) can the following MIPS-like instruction raise? Explain the cause of each exception.

```
# load word from $0(contains zeros) offset 0x120 into register $3
lw  $3,0x0120($0)
```

TLB Handling

In this part of the assignment, you will modify OS/161 to handle TLB faults. Additionally, you need to guarantee that the TLB state is initialized properly on a context switch. One implementation alternative is to invalidate all the TLB entries on a context switch. The entries are then re-loaded by taking TLB faults as pages are referenced. If you do this, be sure to copy any relevant state maintained by the TLB entries back into the page table before invalidating them.

(For example, in order for the paging algorithm to know which pages must be written to disk before eviction, you must make sure that the information about whether a page is dirty or not is properly propagated back into the page table.) An alternative to invalidating everything is to use the 6-bit address space IDs and maintain separate processes in the TLB simultaneously.

We recommend that you separate implementation of the TLB entry replacement algorithm from the actual piece of code that handles the replacement. This will make it easy to experiment with different replacement algorithms if you wish to do so. Refer to the kernel config file section of Assignment 2 on how to add configuration options for TLB replacement policies.

Paging

In this part of the assignment, you will modify OS/161 to handle page faults. When you have completed this problem, your system will generate an exception when a process tries to access an address that is not memory-resident and then handle that exception and continue running the user process.

You will need routines to move a page from disk to memory and from memory to disk.

You will need to decide how to implement backing store (the place on disk where you store virtual pages not currently stored in physical memory). The default `sys161.conf` includes two disks; you can use one of those disks for swapping. Please do swap to a disk and not somewhere else (such as a file). Also, be sure not to use that disk for anything else! To help prevent errors or misunderstandings, please have your system print the location of the swap space when it boots.

You will need to store evicted pages and find them when you need them. You should maintain a bitmap that describes the space in your swap area. Think of the swap area as a collection of chunks, where each chunk holds a page. Use the bitmap to keep track of which chunks are full and which are empty. The empty chunks can be evicted into. You also need to keep track, for each page of a given address space, of which chunk in the swap area it maps onto. When there are too many pages to fit in physical memory, you can write (modified) pages out to swap.

When the time comes to bring a page into memory, you will need to know which physical pages are currently in use. One way to manage physical memory is to maintain a **core map**, a sort of reverse page table. Instead of being indexed by virtual addresses, a core map is indexed by its physical page number and contains the virtual address and address space identifier for the virtual page currently backed by the page in physical memory. When you need to evict a page, you look up the physical address in the core map, locate the address space whose page you are evicting and modify the corresponding state information to indicate that the page will no longer be in memory. Then you can evict the page. If the page is dirty, it must first be written to the backing store. In some systems, the writing of dirty pages to backing store is done in the background. As a result, when the time comes to evict a page, the page itself usually clean (that is, it has been written to backing store, but not modified since then). You could implement this mechanism in OS/161 by creating a thread that periodically examines pages in memory and writes them to backing

store if they are dirty.

Your paging system will also need to support page allocation requests generated by `kmalloc()`. You should review `kmalloc()` to understand how these requests are generated, so that your system will respond to them correctly.

You should implement at least two page-replacement policies, so add kernel configuration options to do so and build your paging system in such a way that it is easy to use different replacement algorithms.

Testing malloc() and free()

This year we have provided you with a malloc implementation. You can grab the source file <u>here</u>. You'll need to add that file to `lib/libc`. (Don't forget to "git add".) Modify `include/stdlib.h` to add function declarations for `malloc()` and `free()`, and add the line `SRCS+=malloc.c` to `lib/libc/Makefile`. For the purposes of these problems, looking at `malloc.c` should be enough.

Now that OS/161 has paging, you can support applications with larger address spaces. The `malloc()` and `free()` library functions are provided in the standard C library. Read the code and answer the following questions in a file called `malloc.txt`:

- **Question 1.** Consider the following (useless) program:

```
/* This is bad code: it doesn't do any error-checking */
#define SIZE 10
#include <stdio.h>
int main (int argc, char **argv) {
        int i;
        void *start, *finish;
        void *res[10];
        start = sbrk(0);
        for (i = 0; i < 10; i++) {
                res[i] = malloc(SIZE);
        }
        finish = sbrk(0);
        /* TWO */
        return 0;
}
```

  1. How many times does the system call `sbrk()` get called from within `malloc()`?
  2. On the i386 platform, what is the numeric value of `(finish - start)`?
- **Question 2.** Suppose that we now insert the following code at location `/* TWO */`:

```
{
        void *x;
        free(res[8]); free(res[7]); free(res[6]);
        free(res[1]); free(res[3]); free(res[2]);
        x = malloc(6*SIZE);  /* MARK */
}
```

  Again on the i386, would `malloc()` call `sbrk()` when doing that last

allocation at the marked line above? What can you say about `x`? Would changing the value of SIZE make a difference (e.g., making it larger)?
- **Question 3.** It is conventional for libc internal functions and variables to be prefaced with "__". Why do you think this is so?
- **Question 4.** The man page for malloc requires that "the pointer returned must be suitably aligned for use with any data type." How does our implementaiton of malloc guarantee this?

Note that the operation of `malloc()` and `free()` is a standard job interview question -- you should understand this code!

You are responsible for making the `malloc()` we give you work; this will involve writing an `sbrk()` system call.

Implementation and Tuning

In this section, we ask you to tune the performance of your virtual memory system. Use the file `performance.txt` as your "lab notebook" for this section. You will undoubtedly want to implement some additional software counters. As a start, we suggest:

- The number of page faults.
- The number of TLB faults.
- The number of page faults that require a synchronous write of a page.

You should add the necessary infrastructure to maintain these statistics as well as any other statistics that you think you will find useful in tuning your system.

Once you have completed all the problems in this assignment and added instrumentation, it is time to tune your operating system.

At a minimum, use the `matmult` and `sort` programs provided to determine your baseline performance (the performance of your system using the default TLB and paging algorithms). Experiment with different TLB and paging algorithms and parameters in an attempt to improve your performance. As before, predict what will happen as you change algorithms and parameters. Compare the measured to the predicted results; obtaining a different result from what you expect might indicate a bug in your understanding, a bug in your implementation, or a bug in your testing. Figure out where the error is! We suggest that you tune the TLB and paging algorithms separately and then together. You may not rewrite the `matmult` program to improve its performance.

You should add other test programs to your collection as you tune performance, otherwise your system might perform will at matrix multiplication and sorting, but little else. Try to introduce programs with some variation in their uses of the memory system.

Provide a complete summary of the performance analysis you conduct. Include a thorough explanation of the performance benefits of your optimizations.

What to hand in

Design Document

Your design document must include:

- Answers to the written problems above.
- How you are going to add paging and TLB handling. Include an integration plan (how your code fits into the existing framework). Be sure to include a description of any new data structures you need.
- A description of the TLB and paging algorithms that you will implement.

Assignment code

- Your revised design document.
- A script of your operating system running the matmult program.
- Your performance.txt file.

After finalizing all of your commits tag your repository `asst3-end`. We will be checking out and grading this version of your repository.

Strategy

The first step is understanding how TLB and page faults occur. To make this task easier, one person in your group should study TLB faults and the other should study page faults.

However, the time required to implement page fault handling is longer than the time required to implement TLB fault handling. You should divide the virtual memory implementation into several small and well-defined modules so that you can both work on it as soon as one of you has completed the TLB implementation. Get together as early as possible to share what you each have discovered.

Look at the code `system161/src/mipseb/mips.c` to see how TLB faults are generated. Then examine the `vm_fault()` handler in `os161/src/kern/arch/mips/mips/dumbvm.c`. What changes must you add to support TLB and page faults?

Some of the key issues are:

- What will your page tables look like?
- What should you put in each PTE (page table entry)?
- In what order can TLB faults and page faults occur? For example, can a page fault occur without causing a TLB fault?

When you have completed your initial implementation of the TLB and virtual memory, one partner can begin experimenting with `matmult` while the other writes other test programs. Stay in touch and test each other's code.

Do your best to break your partner's implementation, then help him/her fix it. When you write test programs, think about verifying that your replacement algorithms are working correctly (i.e., "If I run this program it should generate exactly *n* TLB faults with your algorithm; does it?").

Review each other's designs. Think about how you might have implemented the different parts and compare ideas.

As your system begins to stabilize, begin to work on the performance tuning. Continue to test and fix bugs. Take turns testing and tuning. Be sure to keep a careful log of your performance experiments so you can trade it back and forth.

If you have comments or suggestions, email *p536@cs.indiana.edu*

Author: Andrew Lumsdaine   *Created:* Sat Aug 25, 2001
E-Mail: p536@cs.indiana.edu   *Modified:* Mon 29-Oct-2012
                                                                                                    Copyright ©1997-2013